

# CM2208: Scientific Computing

## Fourier Transform 1:

### Digital Signal and Image Processing

#### Fourier Theory

Prof. David Marshall

School of Computer Science & Informatics

# Fourier Transform

## Moving into the Frequency Domain


The **Frequency domain** can be obtained through the transformation, via **Fourier Transform (FT)**, from

- one (**Temporal (Time)** or **Spatial**) domain
  - to the other
- **Frequency** Domain
  - We do not think in terms of signal or pixel intensities but rather underlying sinusoidal waveforms of varying frequency, amplitude and phase.

## Applications of Fourier Transform

Numerous Applications including:

- Essential tool for Engineers, Physicists, Mathematicians and Computer Scientists
  - Fundamental tool for Digital Signal Processing and Image Processing
  - Many types of Frequency Analysis:
    - Filtering
    - Noise Removal
    - Signal/Image Analysis
    - Simple implementation of Convolution
    - Audio and Image Effects Processing.
    - Signal/Image Restoration — e.g. Deblurring
    - Signal/Image Compression — MPEG (Audio and Video), JPEG use related techniques.
    - Many more .....



## Introducing Frequency Space

## 1D Audio Example

Lets consider a 1D (e.g. Audio) example to see what the different domains mean:

Consider a **complicated sound** such as the a **chord** played on a **piano** or a **guitar**.

We can describe this sound in two related ways:

**Temporal Domain** : Sample the **amplitude** of the sound many times a second, which gives an approximation to the sound as a **function** of **time**.



**Frequency Domain** : Analyse the sound in terms of the **pitches** of the notes, or **frequencies**, which make the sound up, recording the **amplitude** of each frequency.



## Fundamental Frequencies

- D<sub>b</sub> : 554.40Hz  
F : 698.48Hz  
A<sub>b</sub> : 830.64Hz  
C : 1046.56Hz

plus harmonics/partial frequencies ....

# Back to Basics


### An 8 Hz Sine Wave

A signal that consists of a **sinusoidal** wave at **8 Hz**

- 8 Hz means that wave is completing 8 cycles in 1 second
  - The **frequency** of that wave is 8 Hz.

From the **frequency domain** we can see that the composition of our signal is

- **one peak** occurring with a frequency of 8 Hz — there is only one sine wave here.
    - with a **magnitude/fraction** of **1.0** i.e. it is the **whole signal**.



## 2D Image Example

## What do Frequencies in an Image Mean?

Now images are no more complex really:

- Brightness along a line can be recorded as a set of values measured at equally spaced distances apart,
  - Or equivalently, at a set of spatial frequency values.
  - Each of these frequency values is a frequency component.
  - An image is a 2D array of pixel measurements.
  - We form a 2D grid of spatial frequencies.
    - A given frequency component now specifies what contribution is made by data which is changing with specified x and y direction spatial frequencies.

## Frequency components of an image


### What do Frequencies in an Image Mean? (Cont.)

- Large values at **high** frequency components then the data is changing rapidly on a short distance scale.
    - *e.g.* a **page of text**
    - **However**, **Noise** contributes (very) **High Frequencies** also
  - Large **low** frequency components then the large scale features of the picture are more important.  
*e.g.* a single fairly simple object which occupies most of the image.

# Visualising Frequency Domain Transforms

## Sinusoidal Decomposition

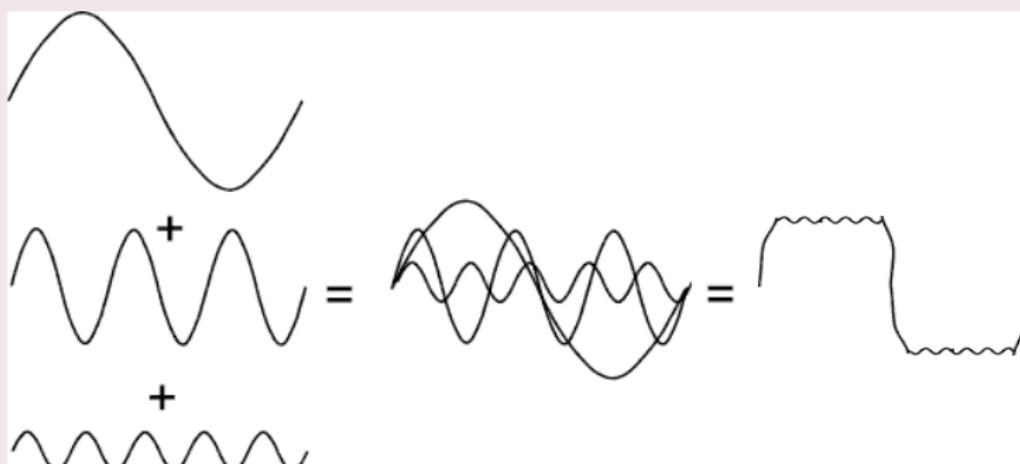
- Any digital signal (function) can be decomposed into purely sinusoidal components
  - Sine waves of different size/shape — varying amplitude, frequency and phase.
- When added back together they reconstitute the original signal.
- The Fourier transform is the tool that performs such an operation.



# Summing Sine Waves. Example: to give a Square(ish) Wave

Digital signals are composite signals made up of many sinusoidal frequencies

- A 200Hz digital signal (**square(ish) wave**) may be a composed of 200, 600, 1000, etc. sinusoidal signals which sum to give:



# Summary so far

## So What Does All This Mean?

Transforming a signal into the frequency domain allows us

- To see what sine waves make up our underlying signal
- E.g.
  - One part sinusoidal wave at 50 Hz and
  - Second part sinusoidal wave at 200 Hz.
  - *Etc.*
- More **complex** signals will give more complex decompositions but the idea is exactly the same.

# How is this Useful then?

## Basic Idea of Filtering in Frequency Space

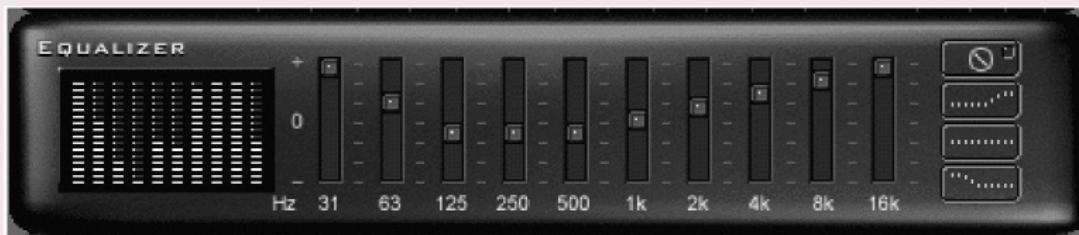
Filtering now involves **attenuating** or **removing** certain frequencies — **easily performed**:

- **Low pass filter** —
  - **Ignore high frequency** noise components — make **zero** or a **very low value**.
  - Only store lower frequency components
- **High Pass Filter** — **opposite of above**
- **Bandpass Filter** — only **allow** frequencies in a **certain range**.

# Visualising the Frequency Domain

## Think Graphic Equaliser


An easy way to visualise what is happening is to think of a graphic equaliser on a stereo system (or some software audio players, e.g. *iTunes*).



# So are we ready for the Fourier Transform?


We have all the Tools....

- This lecture, so far, (hopefully) set the context for Frequency decomposition.
- Past Maths Lectures:
  - Odd/Even Functions:  $\sin(-x) = -\sin(x)$ ,  $\cos(-x) = \cos(x)$
  - Complex Numbers: Phasor Form  $re^{i\phi} = r(\cos\phi + i\sin\phi)$
  - Calculus Integration:  $\int e^{kx}dx = \frac{e^{kx}}{k}$
- Digital Signal Processing:
  - Basic Waveform Theory. Sine Wave  $y = A \cdot \sin(2\pi \cdot n \cdot F_w / F_s)$   
where:  $A$  = amplitude,  $F_w$  = wave frequency,  $F_s$  = sample frequency,  $n$  is the sample index.
  - Relationship between Amplitude, Frequency and Phase:



- Cosine is a Sine wave  $90^\circ$  out of phase
- Impulse Responses
- DSP + Image Proc.: Filters and other processing, Convolution

# Snapshots at jasonlove.com



Professor Herman stopped when he heard that  
unmistakable thud – another brain had imploded.

# Fourier Theory

## Introducing The Fourier Transform

The tool which **converts** a **spatial** or **temporal** (real space) **description** of **audio/image** data, for example, into one in terms of its **frequency components** is called the **Fourier transform**

The new version is usually referred to as the **Fourier space description** of the data.

We then essentially process the data:

- *E.g.* for **filtering** basically this means attenuating or setting certain frequencies to zero

We then need to **convert data back** (or **invert**) to **real audio/imagery** to use in our applications.

The corresponding **inverse** transformation which turns a Fourier space description back into a real space one is called the **inverse Fourier transform**.

# 1D Fourier Transform

## 1D Case (e.g. Audio Signal)

Considering a **continuous** function  $f(x)$  of a single variable  $x$  representing distance (or time).

The **Fourier transform** of that function is denoted  $F(u)$ , where  $u$  represents **spatial** (or **temporal**) **frequency** is defined by:

$$F(u) = \int_{-\infty}^{\infty} f(x) e^{-2\pi i x u} dx.$$

**Note:** In general  $F(u)$  will be a **complex** quantity *even though* the original data is purely **real**.

- The meaning of this is that not only is the **magnitude** of each **frequency** present important, but that its **phase relationship** is **too**.
- Recall **Phasors** from **Complex Number Lectures**.
  - $e^{-2\pi i x u}$  above is a **Phasor**.

# Inverse Fourier Transform

## Inverse 1D Fourier Transform

The **inverse Fourier transform** for regenerating  $f(x)$  from  $F(u)$  is given by

$$f(x) = \int_{-\infty}^{\infty} F(u)e^{2\pi i xu} du,$$

which is rather similar to the (forward) Fourier transform


- except that the **exponential term has the opposite sign.**
- It is **not negative**

# Fourier Transform Example

## Fourier Transform of a Top Hat Function

Let's see how we compute a Fourier Transform: consider a particular function  $f(x)$  defined as

$$f(x) = \begin{cases} 1 & \text{if } |x| \leq 1 \\ 0 & \text{otherwise,} \end{cases}$$



# The Sinc Function (1)

## We derive the Sinc function

So its Fourier transform is:

$$\begin{aligned} F(u) &= \int_{-\infty}^{\infty} f(x)e^{-2\pi i x u} dx \\ &= \int_{-1}^{1} 1 \times e^{-2\pi i x u} dx \\ &= \frac{-1}{2\pi i u} (e^{2\pi i u} - e^{-2\pi i u}) \end{aligned}$$

Now (refer to [Complex Numbers Lectures](#)/[Maths Formula Sheet](#) Handout)

$$\begin{aligned} \sin \theta &= \frac{e^{i\theta} - e^{-i\theta}}{2i}, \text{ So:} \\ F(u) &= \frac{\sin 2\pi u}{\pi u}. \end{aligned}$$

In this case,  $F(u)$  is **purely real**, which is a consequence of the original data being **symmetric** in  $x$  and  $-x$ .

- $f(x)$  is an **even** function.


A graph of  $F(u)$  is shown overleaf.

This function is often referred to as the **Sinc function**.

# The Sinc Function Graph

## The Sinc Function

The Fourier transform of a top hat function, the **Sinc function**:



# The 2D Fourier Transform

2D Case (e.g. Image data)

If  $f(x, y)$  is a function, for example **intensities** in an **image**, its **Fourier transform** is given by

$$F(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-2\pi i(xu+yv)} dx dy,$$

and the **inverse transform**, as might be expected, is

$$f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(u, v) e^{2\pi i(xu+yv)} du dv.$$

# The Discrete Fourier Transform

But All Our Audio and Image data are Digitised!!

Thus, we need a *discrete* formulation of the Fourier transform:

- Assumes regularly spaced data values, and
- Returns the value of the Fourier transform for a set of values in frequency space which are **equally spaced**.

This is done quite naturally by replacing the integral by a summation, to give the *discrete Fourier transform* or **DFT** for short.

# 1D Discrete Fourier transform (DFT)

## 1D Case:

In 1D it is convenient now to assume that  $x$  goes up in steps of 1, and that there are  $N$  samples, at values of  $x$  from 0 to  $N - 1$ .

So the DFT takes the form

$$F(u) = \frac{1}{N} \sum_{x=0}^{N-1} f(x) e^{-2\pi i x u / N},$$

while the inverse DFT is

$$f(x) = \sum_{u=0}^{N-1} F(u) e^{2\pi i x u / N}.$$

**NOTE:** Minor changes from the continuous case are a factor of  $1/N$  in the exponential terms, and also the factor  $1/N$  in front of the forward transform which does not appear in the inverse transform.

# 2D Discrete Fourier transform

## 2D Case

The **2D DFT** works is similar.

So for an  $N \times M$  grid in  $x$  and  $y$  we have

$$F(\mathbf{u}, \mathbf{v}) = \frac{1}{NM} \sum_{x=0}^{N-1} \sum_{y=0}^{M-1} f(x, y) e^{-2\pi i (x\mathbf{u}/N + y\mathbf{v}/M)},$$

and

$$f(x, y) = \sum_{u=0}^{N-1} \sum_{v=0}^{M-1} F(u, v) e^{2\pi i (xu/N + yv/M)}.$$

# Balancing the 2D DFT

## Most Images are Square

Often  $N = M$ , and it is then more convenient to redefine  $F(u, v)$  by multiplying it by a factor of  $N$ , so that the **forward** and **inverse** transforms are more **symmetric**:

$$F(u, v) = \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) e^{-2\pi i(xu+yv)/N},$$

and

$$f(x, y) = \frac{1}{N} \sum_{u=0}^{N-1} \sum_{v=0}^{N-1} F(u, v) e^{2\pi i(xu+yv)/N}.$$

# Fourier Transforms in MATLAB

## fft() and fft2()

MATLAB provides functions for 1D and 2D **Discrete Fourier Transforms (DFT)**:

**fft(X)** is the 1D discrete Fourier transform (DFT) of **vector X**. For **matrices**, the FFT operation is applied to **each column** — **NOT** a 2D DFT transform.

**fft2(X)** returns the 2D Fourier transform of matrix X. If X is a vector, the result will have the same orientation.

**fftn(X)** returns the N-D discrete Fourier transform of the **N-D array X**.

**Inverse DFT ifft(), ifft2(), ifftn()** perform the **inverse** DFT.

See appropriate MATLAB **help/doc** pages for **full details**.

Plenty of examples to Follow.




See also: **MATLAB Docs Image Processing → User's Guide → Transforms → Fourier Transform**

# Visualising the Fourier Transform

## Visualising the Fourier Transform

Having computed a DFT it might be useful to visualise its result:

- It's useful to visualise the Fourier Transform
- Standard tools
- Easily plotted in MATLAB



# The Magnitude Spectrum of Fourier Transform

Recall that the Fourier Transform of our **real** audio/image data is always **complex**

- **Phasors:** This is how we encode the **phase** of the underlying signal's **Fourier Components**.

How can we visualise a complex data array?

Back to Complex Numbers:

Magnitude spectrum **Compute the absolute value of the complex data:**

$$|F(k)| = \sqrt{F_R^2(k) + F_I^2(k)} \text{ for } k = 0, 1, \dots, N - 1$$

where  $F_R(k)$  is the **real** part and  $F_I(k)$  is the **imaginary** part of the  $N$  sampled Fourier Transform,  $F(k)$ .

**Recall MATLAB:** `Sp = abs(fft(X,N))/N;`  
**(Normalised form)**

# The Phase Spectrum of Fourier Transform

## The Phase Spectrum

### Phase Spectrum

The Fourier Transform also represent phase, the **phase spectrum** is given by:

$$\varphi = \arctan \frac{F_I(k)}{F_R(k)} \text{ for } k = 0, 1, \dots, N - 1$$

Recall MATLAB: `phi = angle(fft(X,N))`

# Relating a Sample Point to a Frequency Point

When **plotting graphs** of *Fourier Spectra* and doing other DFT processing we may wish to **plot** the x-axis in **Hz (Frequency)** rather than **sample point** number  $k = 0, 1, \dots, N - 1$

There is a **simple relation** between the two:

- The sample points go in steps  $k = 0, 1, \dots, N - 1$
- For a given sample point  $k$  the frequency relating to this is given by:

$$f_k = k \frac{f_s}{N}$$

where  $f_s$  is the *sampling frequency* and  $N$  the **number** of samples.

- Thus we have **equidistant frequency steps** of  $\frac{f_s}{N}$  ranging from 0 Hz to  $\frac{N-1}{N} f_s$  Hz

# MATLAB Fourier Frequency Spectra Example

## fourierspectraeg.m

```

N=16;
x=cos(2*pi*2*(0:1:N-1)/N);

figure(1)
subplot(3,1,1);
stem(0:N-1,x,'.');
axis([-0.2 N -1.2 1.2]);
legend('Cosine signal x(n)');
ylabel('a');
xlabel('n \rightarrow');

X=abs(fft(x,N))/N;
subplot(3,1,2);stem(0:N-1,X,'.');
axis([-0.2 N -0.1 1.1]);
legend('Magnitude spectrum |X(k)|');
ylabel('b');
xlabel('k \rightarrow')

N=1024;
x=cos(2*pi*(2*1024/16)*(0:1:N-1)/N);

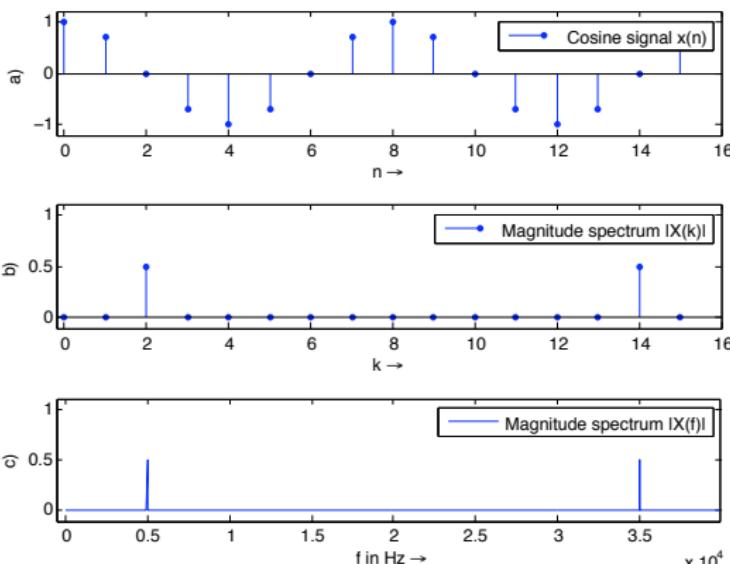
FS=40000;
f=((0:N-1)/N)*FS;
X =abs(fft(x,N))/N;
subplot(3,1,3);plot(f,X);
axis([-0.2*44100/16 max(f) -0.1 1.1]);
legend('Magnitude spectrum |X(f)|');
ylabel('c');
xlabel('f in Hz \rightarrow')

figure(2)
subplot(3,1,1);
plot(f,20*log10(X./(0.5)));
axis([-0.2*44100/16 max(f) ... -45 20]);
legend('Magnitude spectrum |X(f)| ... in dB');
ylabel('|X(f)| in dB \rightarrow');
xlabel('f in Hz \rightarrow')

```


# MATLAB Fourier Frequency Spectra Example Output

fourierspectraeg.m produces the following:



# Magnitude Spectrum in dB

**Note:** It is common to plot both spectra magnitude (also frequency ranges not show here) on a dB/log scale:  
(Last Plot in fourierspectraeg.m)



# Time-Frequency Representation: Spectrogram

## Spectrogram

It is often **useful** to look at the **frequency distribution** over a **short-time**:

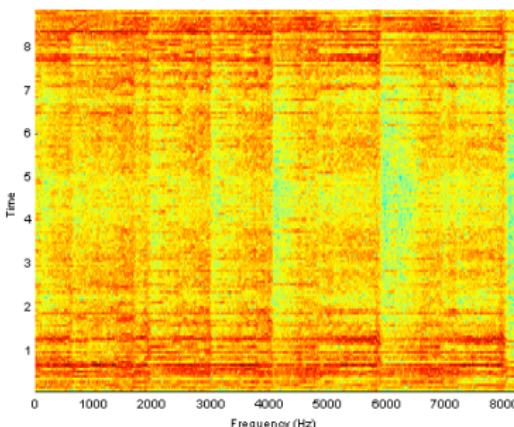
- Split signal into  $N$  segments
- Do a **windowed Fourier Transform** — **Short-Time Fourier Transform (STFT)**
  - Window needed to reduce *leakage* effect of doing a shorter sample SFFT.
  - Apply a **Blackman**, **Hamming** or **Hanning** Window
- MATLAB function does the job: **Spectrogram** — see **help spectrogram**
- See also MATLAB's **specgramdemo**

# MATLAB spectrogram Example

spectrogrammeg.m

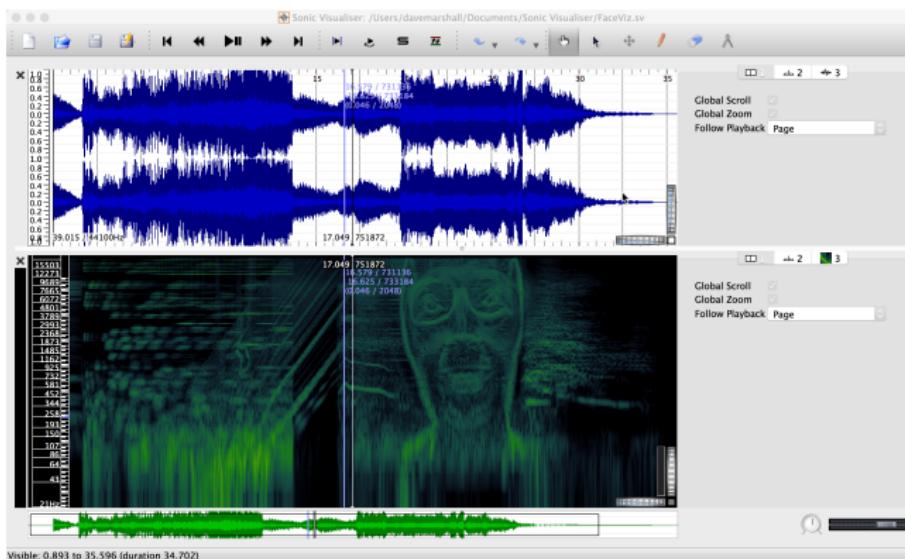
```
load('handel')
[N M] = size(y);
figure(1)
spectrogram(fft(y,N),512,20,1024,Fs);
```

Produces the following:



# Aphex Twin Spectrogram

Aphex Twin famously<sup>1</sup> embedded images in the spectrogram of a few tracks on his Windowlicker EP. His face on Track 2 “Formula” or “Equation” (Full title:  $\Delta M_{i-1} = -\alpha \sum_{n=1}^N D_i[n][\sum_{\sigma \in C[i]} F_{ji}[n-1] + F_{exti}[n-1]]$ !!:

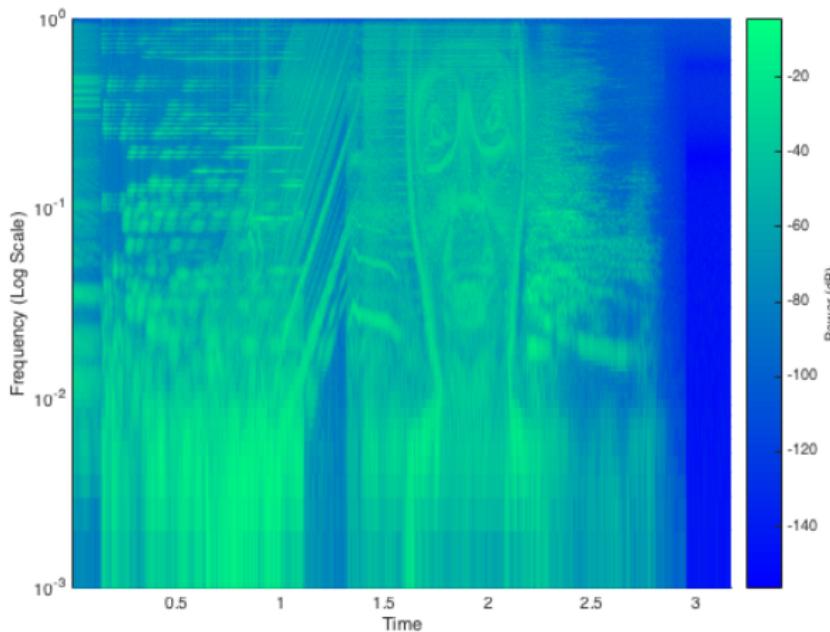


<sup>1</sup>See [here for web link](#) to other examples of embedded image Spectrograms

# Matlab Code to show the Aphex Twin Spectrogram

Previous slide use the free and excellent [Sonic Visualiser](#)

We of course know how to display the image in MATLAB:



# Matlab Code to show the Aphex Twin Spectrogram

## Aphex\_Spectrogram.m

```
aphex = audioread('FormulaSnippet.wav');

mono = (aphex(:,1) + aphex(:,2))/2;

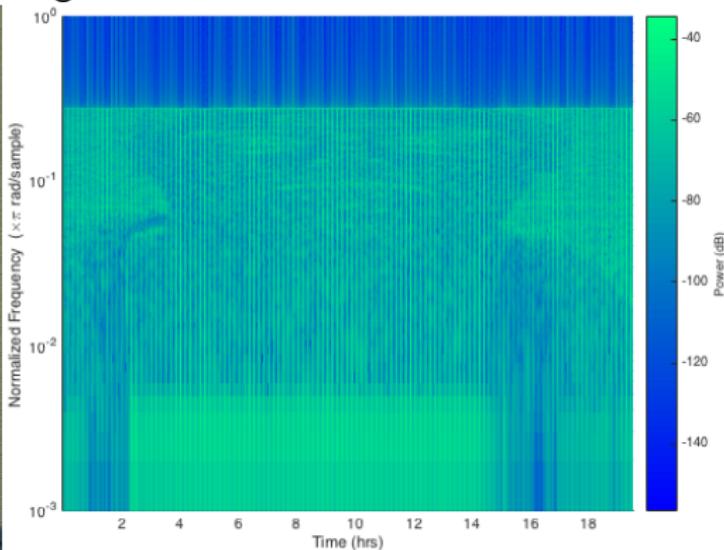
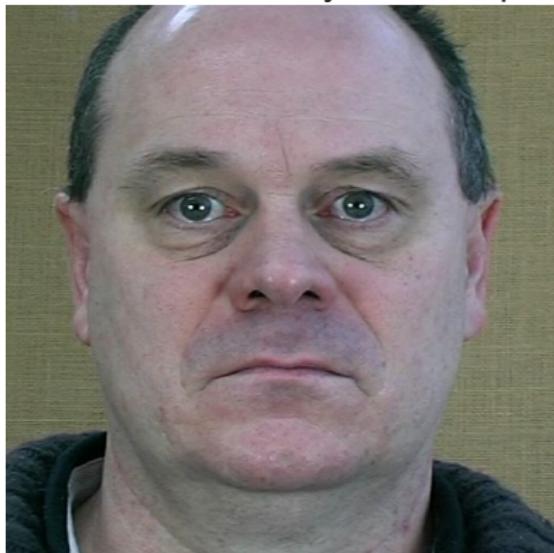
spectrogram(mono,1024,120,2048,'power','yaxis');
set(gca,'YScale','log');
colormap('winter');
xlabel('Time')
ylabel('Frequency (Log Scale)')
```

**Note:** we change the display of the spectrogram to a **log scale**, which looks better.

**Audio clip here:** [FormulaSnippet.wav](#)

# So what does my face sound like?

Let's embed my face in spectrogram:



It sounds like this:



# Image to Sound Conversion

## Daphex.m

```
figure(1);
imshow(imread('Dave_Frame0001.jpg'));

dave_im2snd = im2sound('Dave_Frame0001', 'jpg', 44100, 40,6000,0.00002, 10);

sound(dave_im2snd,44100);

figure(2);

spectrogram(dave_im2snd,1024,120,2048,'power','yaxis');
set(gca,'YScale','log');

colormap('winter');

shg;
```

Image used here: [Dave\\_Frame0001.jpg](#)

# Image to Sound Conversion<sup>2</sup>

## im2sound.m (Usage)

```
function [final_sound] = im2sound(filename, ext, f_sample, f_low, ...
f_high, amp_mod, sample_t)

%INPUTS:
%'filename' - Name of the image to be encoded (not including extension)
%'ext' - Extension of the image (not including "." at the beginning).
%'f_sample' - Sampling frequency (Hz)
%'f_low' - Lowest frequency (Hz) (e.g. 40)
%'f_high' - Highest frequency (Hz) (e.g. 6000)
%'amp_mod' - Multiplication factor for the amplitude. Decrease until
%image is clear. Too high and the waveform clips. Too low and the image
%is very dark (e.g. 0.00002)
%'sample_t' - Duration of the sample in seconds. Longer samples have
%better quality (e.g. 10)

%OUTPUTS:
%'final_sound' - the final sound containing the image. This is
%automatically saved to a .wav file with the original image filename
```

<sup>2</sup>Original Code from [MATLAB Central](#)

# Image to Sound Conversion

## im2sound.m (Code)

```
function [final_sound] = im2sound(filename, ext, f_sample, f_low, ...
    f_high, amp_mod, sample_t)

%%%%

%INITIALISING VARIABLES:
%The waveform at each time point. This is reset at the beginning of each
%time point
temp_sound = 0;
%The final waveform
final_sound = 0;

%MAIN BODY
%Loading the sample image and calculating the image size
raw_im = imread(strcat(filename,'.',ext));
size_raw_im = size(raw_im);

%Making a frequency table for the height of the image. Each row of the
%image is assigned a particular frequency from the corresponding row of
%this table. The frequencies are linearly distributed between the highest and
%lowest user-defined frequencies. "f_step" is the increment between each
%adjacent frequency
f_step = (f_high - f_low)/size_raw_im(1,1);
f_table = (f_high:-f_step:f_low);
```

# Image to Sound Conversion

## im2sound.m (Code)

```
%The final sound will dwell on each column of the image for a specific
%time. This time is defined by "t_start" and "t_end". It depends on how
%long the user determined the sound-clip should be and how wide (how many
%columns) the image is.
t_step = (sample_t/size_raw_im(1,2));

%Initial values for the start and end times. These will be increased at
%the end of each loop iteration (when the script moves onto the next column
%of the image).
t_start = 0;
t_end = t_step;

%The loop which generates the sound file. At each iteration it generates a
%segment of the final sound file, which is temporarily saved to
%"temp_sound". This segment is built up of frequencies from that
%particular column of the image.
for j = 1:size_raw_im(1,2)
    %Initialising the variable (the sound for each frequency (row) is added
    %to the existing sound)
    temp_sound = 0;

    %Setting the time in matrix format
    t = t_start:1/f_sample:(t_end);
```

# Image to Sound Conversion

## im2sound.m (Code)

```
%For each iteration of this loop, the script goes down the current
%column of the image and generates a waveform of the frequency
%specified in "f_table". The amplitude of the waveform is determined
%by the pixel intensity. This generated waveform is added to all the
%previously generated waveforms in that particular column
for i = size_raw_im(1,1):-1:1
    temp_sound = temp_sound+ sin(2*pi*t*f_table(i))*...
        double(raw_im(i,j))*amp_mod;

end

%At the end of each column the segment of sound generated is added to
%the end of the existing sound file ("final_sound").
final_sound = cat(2,final_sound,temp_sound);

%The temporary sound is cleared ready for the start of the next column
clear temp_sound

%Moving to the next time frame
t_start = t_start + t_step;
t_end = t_end + t_step;

end

%This saves "final_sound" to the '.wav' file of the same name as the input
%file
wavwrite(final_sound, f_sample, strcat(filename, '.wav'));
```

# Properties of Fourier Transforms

Here are just a few of the other properties of Fourier transforms which are useful when reasoning, or computing, with Fourier transforms:

- Linear Operator
- Shifting
- Scaling
- Rotation
- Zeroth component
- Convolution Theorem (**see next section applications**)

# Fourier Transform: Linear Operator


The Fourier transform is a *linear operator*. This means that

## Theorem


If  $f(x)$  and  $g(x)$  are two functions with Fourier transforms  $F(u)$  and  $G(u)$ , then the Fourier transform of  $af(x) + bg(x)$  where  $a$  and  $b$  are constants is simply  $aF(u) + bG(u)$ .

## Linear Operator Simple Example


**Image 1**






**Image 2**



**Image 3**



## Linear Operator Simple Example



# Linear Operator Simple Example

## Fourier Transform Linear Operator Demo: FFT\_Linear.m

```
% Create a white box on a black background image
M= 256; N = 256;
image1 = zeros(M,N);
box = ones(64,64);
%box at centre
image1(97:160,97:160) = box;

figure(1)
imshow(image1)
title('Image 1');

% Create another white box on a black background image
M= 256; N = 256;
image2 = zeros(M,N);
box = ones(32,32);
%box at centre
image2(37:68,37:68) = box;

figure(2)
imshow(image2)
title('Image 2');
```

# Linear Operator Simple Example

## Fourier Transform Linear Operator Demo: FFT\_Linear.m

```
% Make composite image
image3 = image1 + image2;

figure(3)
imshow(image3)
title('Image 3');

% Compute Fourier Transforms.

imft1 = fft2(double(image1));
imft2 = fft2(double(image2));
imft3 = fft2(double(image3));

figure(4)
imagesc(abs(imft1 + imft2))
title('Image 1 + Image 2 Spectra');

figure(5)
imagesc(abs(imft3))
title('imft3 Spectra');

figure(6)
imagesc(abs(imft1 + imft2 - imft3))
title('Difference (Image 1 + Image 2) - imft3 Spectra');

figure(7)
imshow(fft2(imft1 + imft2));
title('Inverse FT of imft1 + imft2');
```

# Fourier Transform: Shifting

Shifting the real space data through a fixed distance has the effect that

## Theorem

If  $f(x)$  is a function with Fourier transform  $F(u)$ , then the Fourier transform of  $f(x - x_0)$  is given by  $e^{-2\pi i x_0 u} F(u)$ .

- **Note:** that the exponential term here has **unit** magnitude for all values of  $u$ .
  - Thus, the **magnitude** of the **resulting** Fourier transform is left **unchanged**, although the **relative** sizes of the real and imaginary parts are **altered**.

# Shifting Operator Simple Example

Image 1

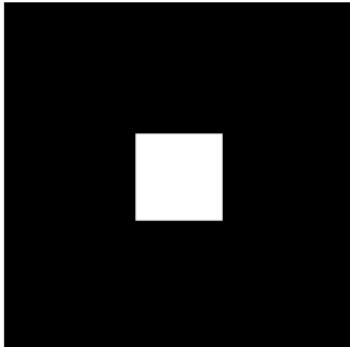
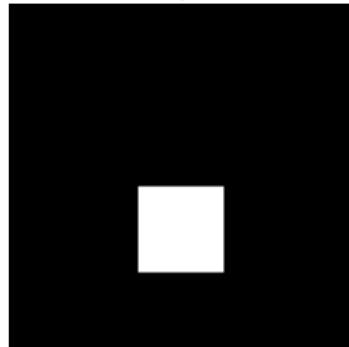
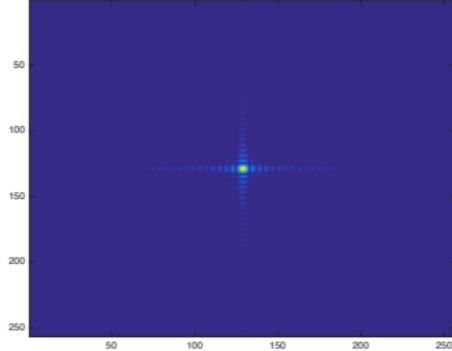



Image 2



FT of Image 1 (imft1)



FT of Image 2 (imft2)



# Shifting Operator Simple Example

Inverse FT of imft3

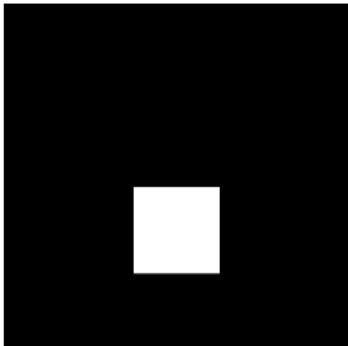
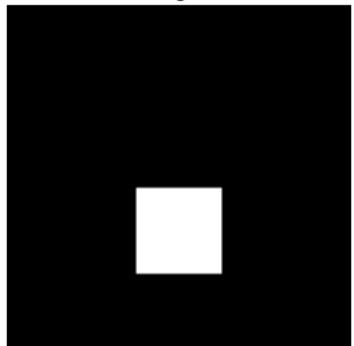
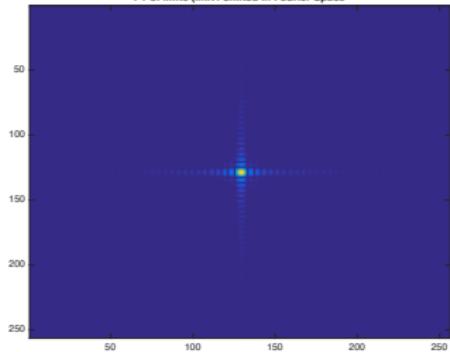


Image 2

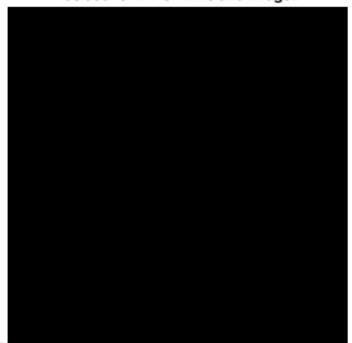


Both Images are the same.

FT of imft3 (imft1 shifted in Fourier Space)



Residual of IFT of imft3 and image2



# Shifting Operator Simple Example

## Fourier Transform Shifting Operator Demo: FFT\_Shift.m

```
% Define Initial box
x0 = 97;
y0 = 97;
xwidth = 64;
ywidth = 64 ;

% Create a white box on a black background image
M= 256; N = 256;
image1 = zeros(M,N);
box = ones(xwidth,ywidth);
%box at centre
image1(x0:x0+xwidth-1,y0:y0+ywidth-1) = box;

figure(1)
imshow(image1)
title('Image 1');

% Define shift
dx= 40;
dy= 0;

% Shift Image 1
image2 = zeros(M,N);
image2(x0+dx:x0+dx+xwidth-1,y0+dy:y0+dy+ywidth-1) = box;

figure(2)
imshow(image2)
title('Image 2');
```

# Shifting Operator Simple Example

## Fourier Transform Shifting Operator Demo: FFT\_Shift.m

```
% Compute Fourier Transforms.  
imft1 = fft2(double(image1));  
imft2 = fft2(double(image2));  
  
figure(3)  
imagesc(abs(fftshift(imft1)));  
title('FT of Image 1 (imft1)');  
  
figure(4)  
imagesc(abs(fftshift(imft2)));  
title('FT of Image 2 (imft2)');  
  
% Define shift in frequency domain  
[yF,xF] = meshgrid(-M/2:M/2-1,-N/2:N/2-1);  
  
% Perform the shift  
imft3=imft1.*exp(-1i*2*pi.*(xF*dx+yF*dy)/256);  
  
figure(5)  
imagesc(abs(fftshift(imft3)));  
title('FT of imft3 (imft1 shifted in Fourier Space');  
  
figure(6)  
imshow(ifft2(imft3));  
title('Inverse FT of imft3');  
  
figure(7)  
imshow(ifft2(imft3)-image2);  
title('Residual of IFT of imft3 and image2');
```

# Fourier Transform: Scaling

If we scale the spacing of the real space data in distance, we have that

## Theorem

*If  $f(x)$  is a function with Fourier transform  $F(u)$ , then the Fourier transform of  $f(ax)$  where  $a$  is a real constant is given by  $\frac{1}{|a|} F(\frac{u}{a})$ .*

## In other words:

- if we **spread out** the data in **real space**, the data is **compressed** into more **compact** region of **Fourier space**.
- This is **intuitive** — if we **double** the spacing of a **grid pattern**, its **spatial frequency** is **halved**.
- **Note** that the **magnitude** of the Fourier representation is also **affected**.

# Scaling Operator Simple Example

Image 1

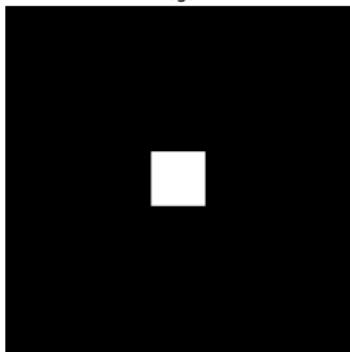
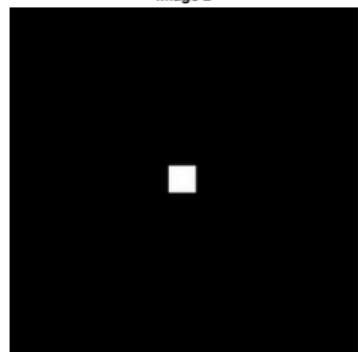
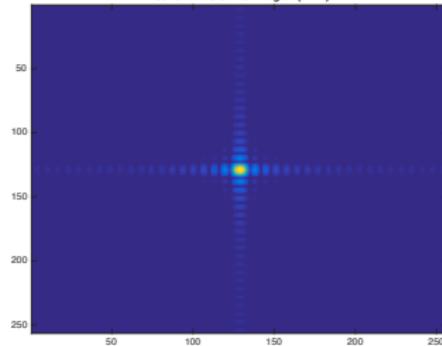



Image 2



Fourier Transform of Image 1 (imft1)



Fourier Transform of Image 2 (imft2)



# Scaling Operator Simple Example

## Fourier Transform Scaling Operator Demo: FFT\_Scale.m

```
% Define Initial box
xwidth = 40;
ywidth = 40;

% Create a white box on a black background image
M= 256; N = 256;

x0 = M/2 - xwidth/2;
y0 = N/2 - ywidth/2;

image1 = zeros(M,N);
box = ones(xwidth,ywidth);
%box at centre
image1(x0:x0+xwidth-1,y0:y0+ywidth-1) = box;

% Define scale
scalex = 0.5;
scaley = 0.5;

% Image 2 = Scaled Image 1
image2 = zeros(M,N);
[px,qx] = rat(scalex);
[py,qy] = rat(scaley);
```

# Scaling Operator Simple Example

## Fourier Transform Scaling Operator Demo: FFT\_Scale.m

```
in_image2 = resample(resample(double(image1),px,qx)',py,qy)';

xrange = int16(M/2-M*scalex/2:M/2+M*scalex/2-1);
yrange = int16(N/2-N*scaley/2:N/2+N*scaley/2-1);
image2(xrange,yrange)= in_image2;

figure(1)
imshow(image1);
title('Image 1');

figure(2)
imshow(image2);
title('Image 2');

% Compute Fourier Transfoms.
imft1 = fft2(double(image1));
imft2 = fft2(double(image2));

figure(3)
imagesc(abs(fftshift(imft1)))
title('Fourier Transform of Image 1 (imft1)');

figure(4)
imagesc(abs(fftshift(imft2)))
title('Fourier Transform of Image 2 (imft2)');
```

# Fourier Transform: Rotation

One useful property in two dimensions is that:

- if we **rotate** the **real space** data, its **Fourier transform** is **rotated** by the **same angle**, or more exactly

## Theorem

If  $f(x, y)$  is a function with Fourier transform  $F(u, v)$ , on expressing these functions in terms of polar coordinates  $r, \theta, \rho, \phi$  where  $x = r \cos \theta, y = r \sin \theta, u = \rho \cos \phi, v = \rho \sin \phi$  so that  $f(x, y)$  and  $F(u, v)$  become  $f(r, \theta)$  and  $F(\rho, \phi)$  respectively, the Fourier transform of  $f(r, \theta + \omega)$  where  $\omega$  is a constant is given by  $F(\rho, \phi + \omega)$ .

- Basically this means the **two-dimensional** Fourier transform is an intrinsic property of the data which is **independent** of our choice of axis directions.

# Rotation Operator Simple Example

Image 1

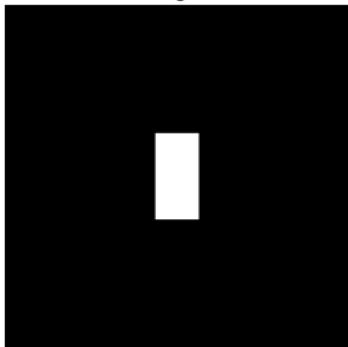
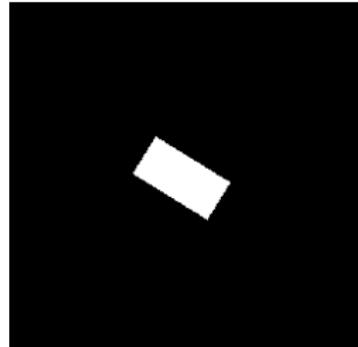
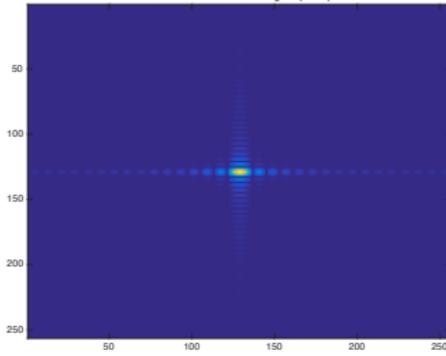


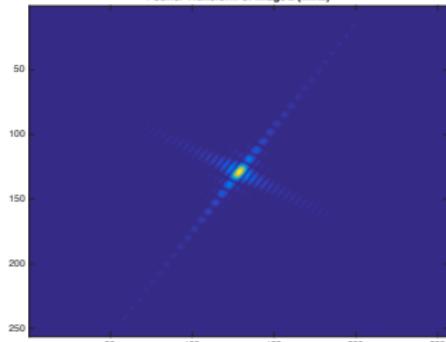
Image 2



Fourier Transform of Image 1 (imft1)



Fourier Transform of Image 2 (imft2)



# Rotation Operator Simple Example

## Fourier Transform Rotate Operator Demo: FFT\_Rotate.m

```
% Define Initial box
xwidth = 64;
ywidth = 32;

% Create a white box on a black background image
M= 256; N = 256;
image1 = zeros(M,N);

x0 = M/2 - xwidth/2;
y0 = N/2 - ywidth/2;

box = ones(xwidth,ywidth);
%box at centre
image1(x0:x0+xwidth-1,y0:y0+ywidth-1) = box;

% Define Rotation in degrees
rot = 45;

image2 = imrotate(image1, rot,'bilinear','crop');

figure(1); imshow(image1);
title('Image 1');

figure(2); imshow(image2);
title('Image 2');

% Compute Fourier Transfoms.
imft1 = fftshift(fft2(double(image1)));
imft2 = fftshift(fft2(double(image2)));

figure(3); imagesc(abs(imft1))
title('Fourier Transform of
Image 1 (imft1)');

figure(4); imagesc(abs(imft2))
title('Fourier Transform of I
mage 2 (imft2)');
```

# Fourier Transform: Zeroth Component

A final property is that the **zeroth component** of the **Fourier space representation** is just the **average data value** (apart from a factor of  $1/N$ , in two dimensions – assuming balanced FT)).  
This is illustrated below for the two-dimensional DFT case:

$$F(u, v) = \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) e^{-2\pi i(xu+yv)/N}, \quad (1)$$

so

$$F(0, 0) = \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) = N\bar{f}(x, y). \quad (2)$$

# Simple Zeroth Component Example

Illustration of zeroth element,  $F(1)$  (1D) or  $F(1,1)$  (2D), being the mean of the data plus a factor  $N$ : FFT\_zeroth.m

```
% 1D Example
y = 1:100
yft = fft(y);
yft(1)/100 % Zeroth Element of y
mean(y) % Mean of Data

% 2D Example
im = imread('cameraman.tif');

[N M] = size(im);

imft = fft2(im);


imft(1,1)/(N*M) % Zeroth Element of imft
mean(mean(im)) % Mean of Data
```

FT is not balanced in MATLAB

# Corollary: Shifting the Fourier Transform, `fftshift()`

## Centring the Frequency of a Fourier Transform

- Most computations of FFT represent the frequency from  $0 — N - 1$  samples (similarly in 2D, 3D etc.) with corresponding frequencies ordered accordingly — the **0** frequency is not really the **centre**.
- We frequently like to visualise the FFT as the **centre of the spectrum**.
- In 1D (Audio/Vector): **swaps the left and right halves of the vector**
- Similarly in 2D (Image/Matrix) we swap the first quadrant with the third and the second quadrant with the fourth:



# The fftshift() MATLAB Command

```
help fftshift()
```

**Y = fftshift(X)** rearranges the outputs of **fft**, **fft2**, and **fftn** by **moving** the zero-frequency component to the **center of the array**.

It is useful for **visualising** a Fourier transform with the zero-frequency component in the **middle** of the spectrum.

For **vectors**, **fftshift(X)** **swaps** the **left** and **right** halves of X.

For **matrices**, **fftshift(X)** swaps the **first** quadrant with the **third** and the **second** quadrant with the **fourth**.